



**BITS Pilani**  
Hyderabad Campus

# OPERATING SYSTEMS (CS F372)

## OS Structures

Dr. Barsha Mitra  
CSIS Dept., BITS Pilani, Hyderabad Campus

# Operating System Services

---

- ❖ **User interface** - almost all operating systems have a user interface (**UI**)
  - ❖ **Command-Line (CLI), Graphics User Interface (GUI)**
- ❖ **Program execution** - system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
- ❖ **I/O operations** - running program may require I/O, which may involve a file or an I/O device
- ❖ **File-system manipulation** - read and write files and directories, create and delete them, search them, list file Information, permission management

# Operating System Services

---

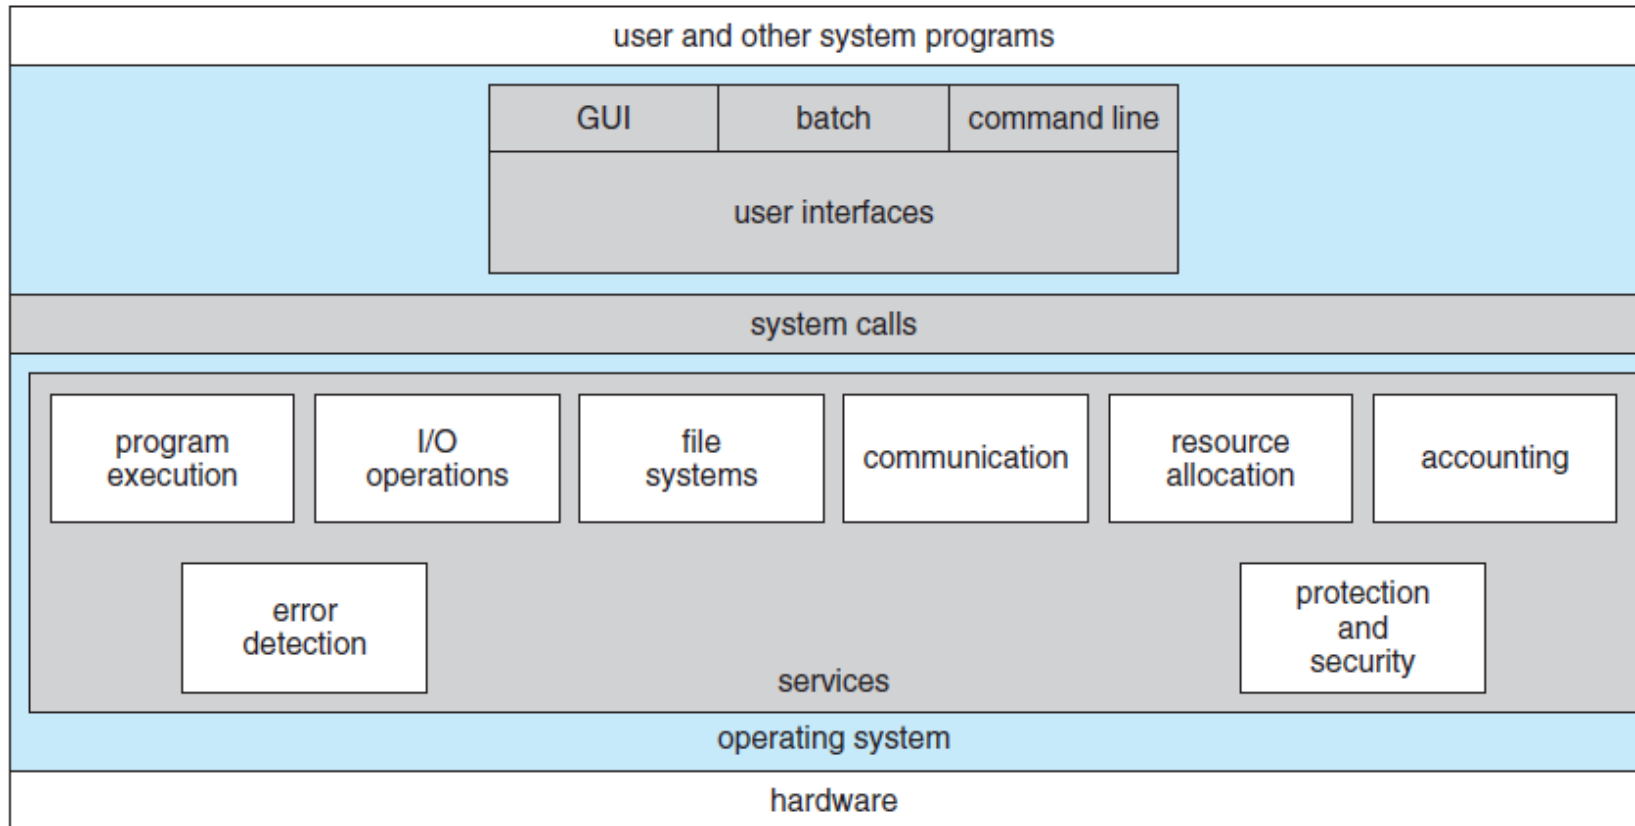
- ❖ **Communications** – processes may exchange information, on the same computer or between computers over a network, **shared memory or message passing**
- ❖ **Error detection** –
  - ❖ OS needs to be constantly aware of possible errors
  - ❖ May occur in CPU and memory h/w, in I/O devices, in user program
  - ❖ OS should take the appropriate action to ensure correct and consistent computing
  - ❖ Take corrective actions

# Operating System Services

---

- ❖ **Resource allocation** – allocating resources like CPU cycles, main memory, file storage, I/O devices for multiple concurrently executing processes
- ❖ **Accounting** – keep track of which users use how much and what kinds of computer resources
- ❖ **Protection and security** –
  - ❖ owners of information stored in a multiuser or networked computer system want to control use of that information
  - ❖ concurrent processes should not interfere with each other or with OS
  - ❖ ensuring that all accesses to system resources is controlled
  - ❖ security of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts

# Operating System Services



# User and Operating-System Interface: CLI

---



- ❖ **CLI** or command interpreter
- ❖ Sometimes implemented in kernel, sometimes by separate program (Unix, Windows)
- ❖ Sometimes multiple flavors implemented – shells
- ❖ Primarily fetches a command from user and executes it

# User and Operating-System Interface: GUI



- ❖ User-friendly interface
- ❖ Usually mouse, keyboard, and monitor
- ❖ Icons represent files, programs, actions, etc.
- ❖ Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a folder))
- ❖ Many systems now include both CLI and GUI interfaces
  - ❖ Microsoft Windows is GUI with CLI “command” shell
  - ❖ Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)



# User and Operating-System Interface: Touchscreen Interface



- ❖ Touchscreen devices require new interfaces
- ❖ Mouse not possible or not desired
- ❖ Actions and selection based on gestures
- ❖ Virtual keyboard for text entry
- ❖ Voice commands





## Choice of Interface

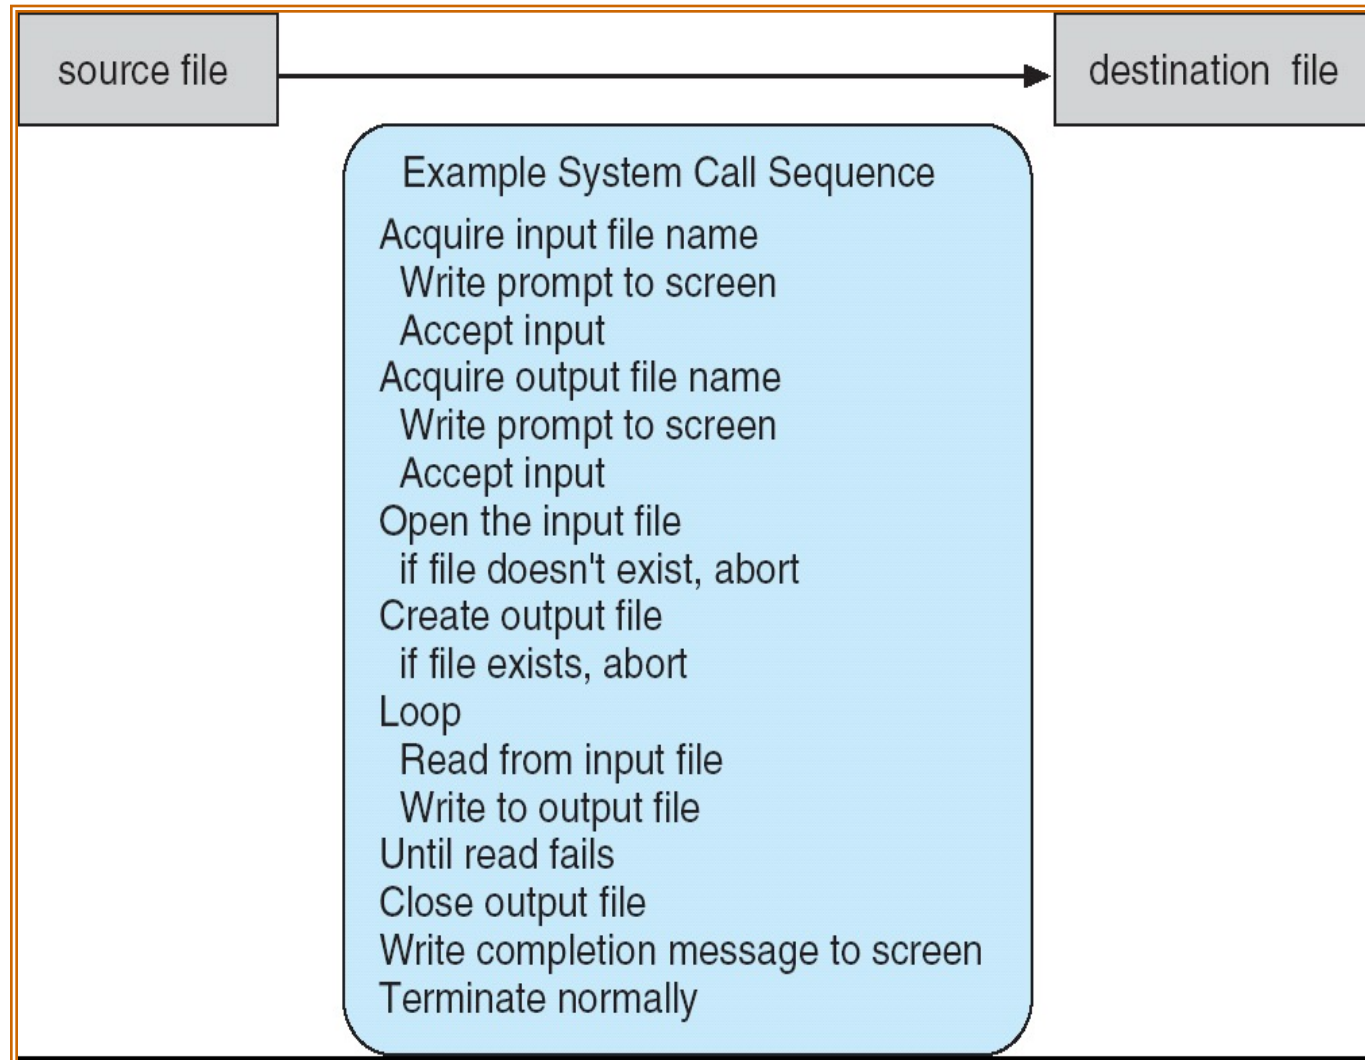


# System Calls

innovate

achieve

lead



# System Calls



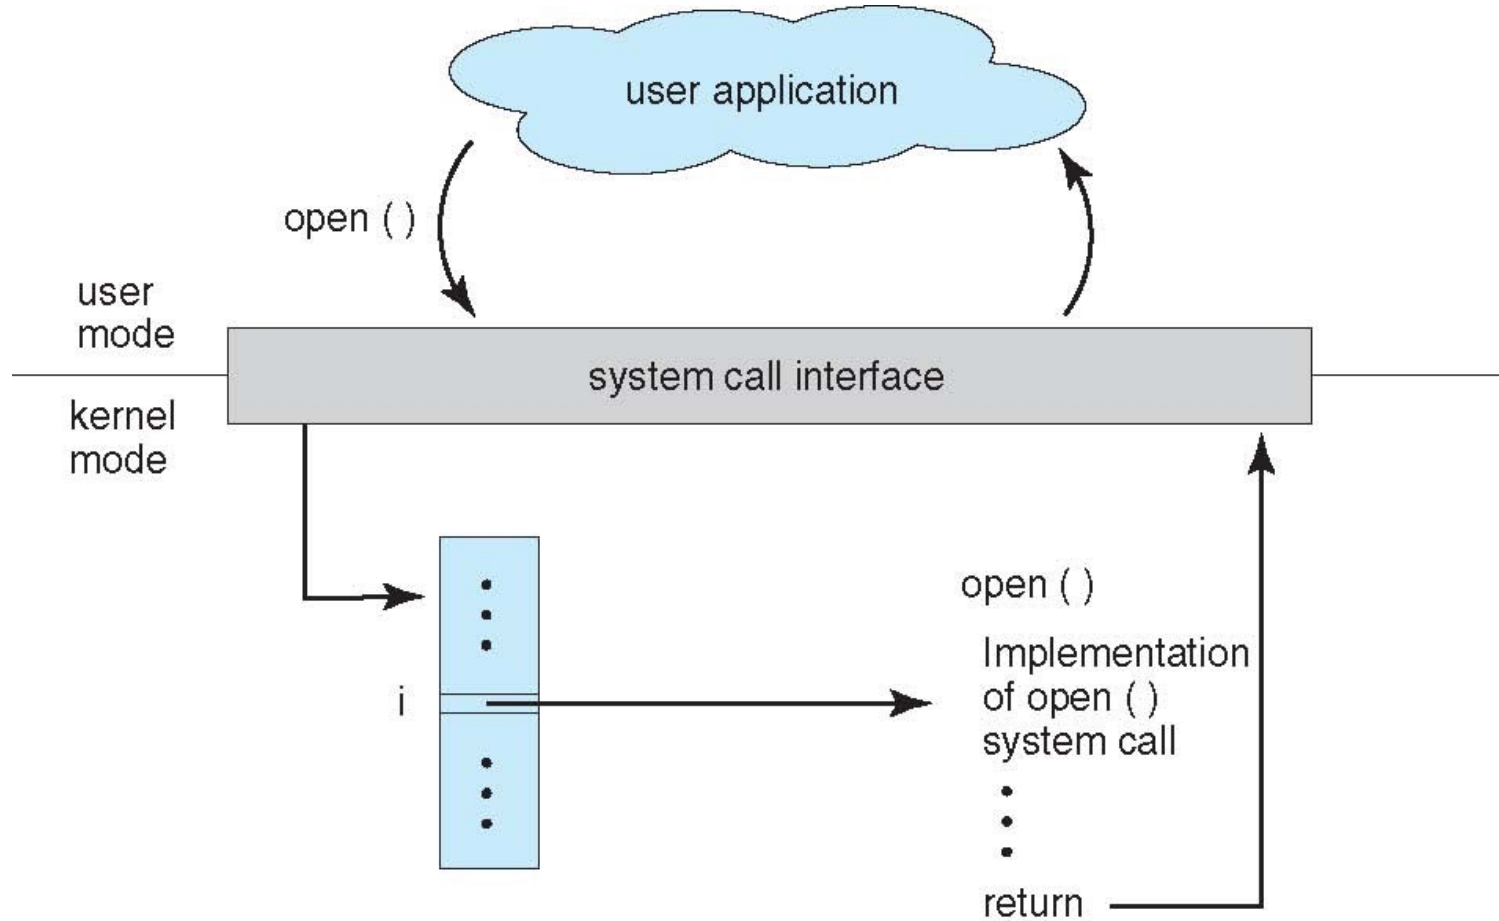
- ❖ Interface to the services provided by the OS
- ❖ Typically written in a high-level language (C or C++)
- ❖ Mostly accessed by programs via a high-level Application Programming Interface (API) rather than direct system call use
- ❖ API specifies a set of functions available to application programmers
- ❖ Programmers access API via code library provided by the OS
- ❖ Three most common APIs are
  - ❖ Win32 API for Windows
  - ❖ POSIX API for POSIX-based systems (including all versions of UNIX, Linux, and Mac OS X)
  - ❖ Java API for the Java virtual machine (JVM)

# System Calls



- ❖ A number is associated with each system call
- ❖ System-call interface maintains a table indexed according to these numbers
- ❖ The system call interface invokes the intended system call in OS kernel and returns status of the system call and any return values
- ❖ The caller need know nothing about how the system call is implemented
- ❖ Just needs to obey API and understand what OS will do as a result of call execution
- ❖ Most details of OS interface hidden from programmer by API
- ❖ Managed by run-time support library (set of functions built into libraries included with compiler)

# System Calls



# Types of System Calls



- **Process control**

- create process, terminate process
- end, abort
- load, execute
- get process attributes, set process attributes

- **File management**

- create file, delete file
- open, close file
- read, write file
- get and set file attributes

- **Device management**

- request device, release device
- read, write
- get device attributes, set device attributes
- logically attach or detach devices

- **Information Maintenance**

- **Communication**

- **Protection**

# Examples of System Calls



	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()



# OS Structure

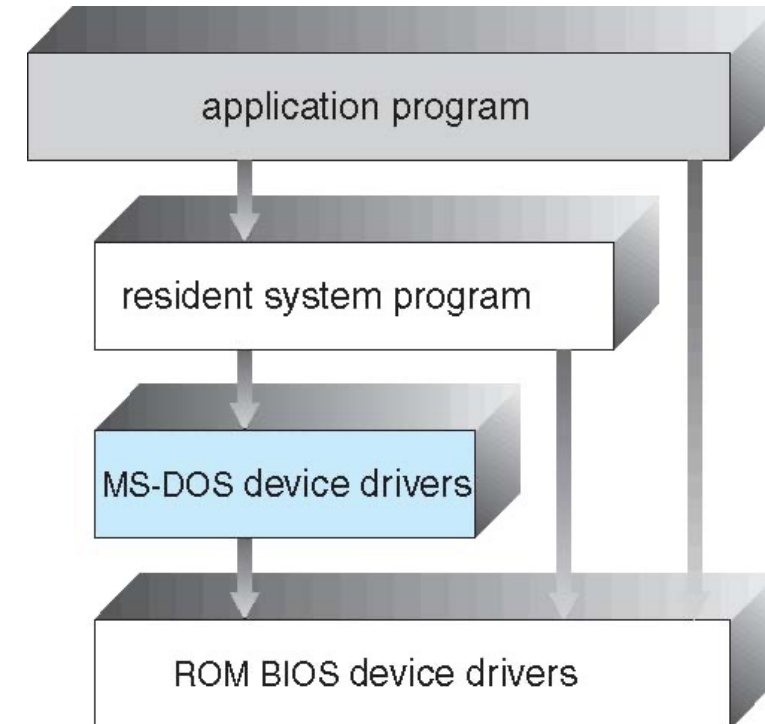


- Simple Structure/ Monolithic Kernel
- Layered Approach
- Microkernels
- Modules
- Hybrid System

# Simple Structure



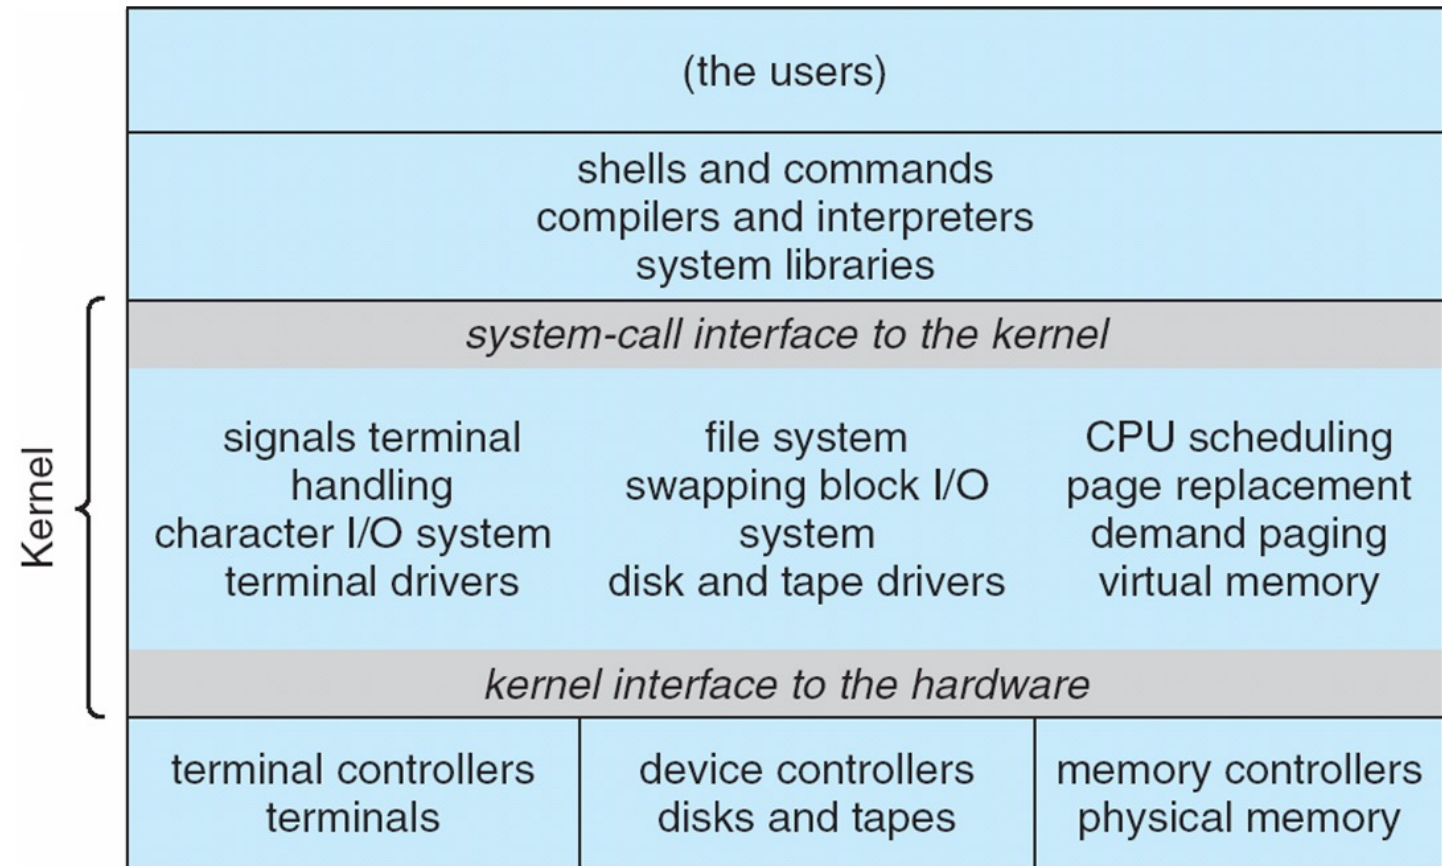
- not divided into modules
- interfaces and levels of functionality are not well separated
- application programs are able to access the basic I/O routines to write directly to the display and disk drives
- vulnerable to malicious programs, causing entire system crashes when user programs fail



# UNIX Architecture



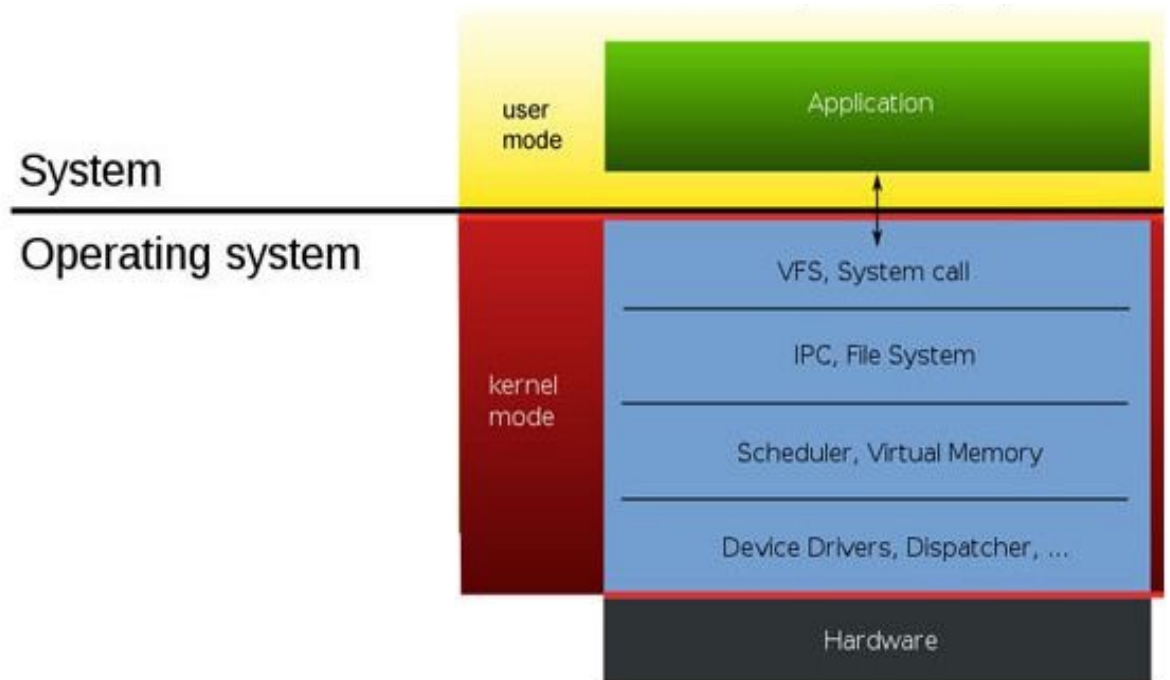
- the original UNIX operating system had limited structuring
- consists of two separable parts
  - Systems programs
  - kernel
    - Consists of everything below the system-call interface and above the physical hardware
    - Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level



# Monolithic Kernel




- entire operating system is working in kernel space
- larger in size
- little overhead in system call interface or in communication within kernel
- faster
- hard to extend
- if a service crashes, whole system is affected
- Eg., - Linux, Solaris, MS-DOS



# Layered Approach



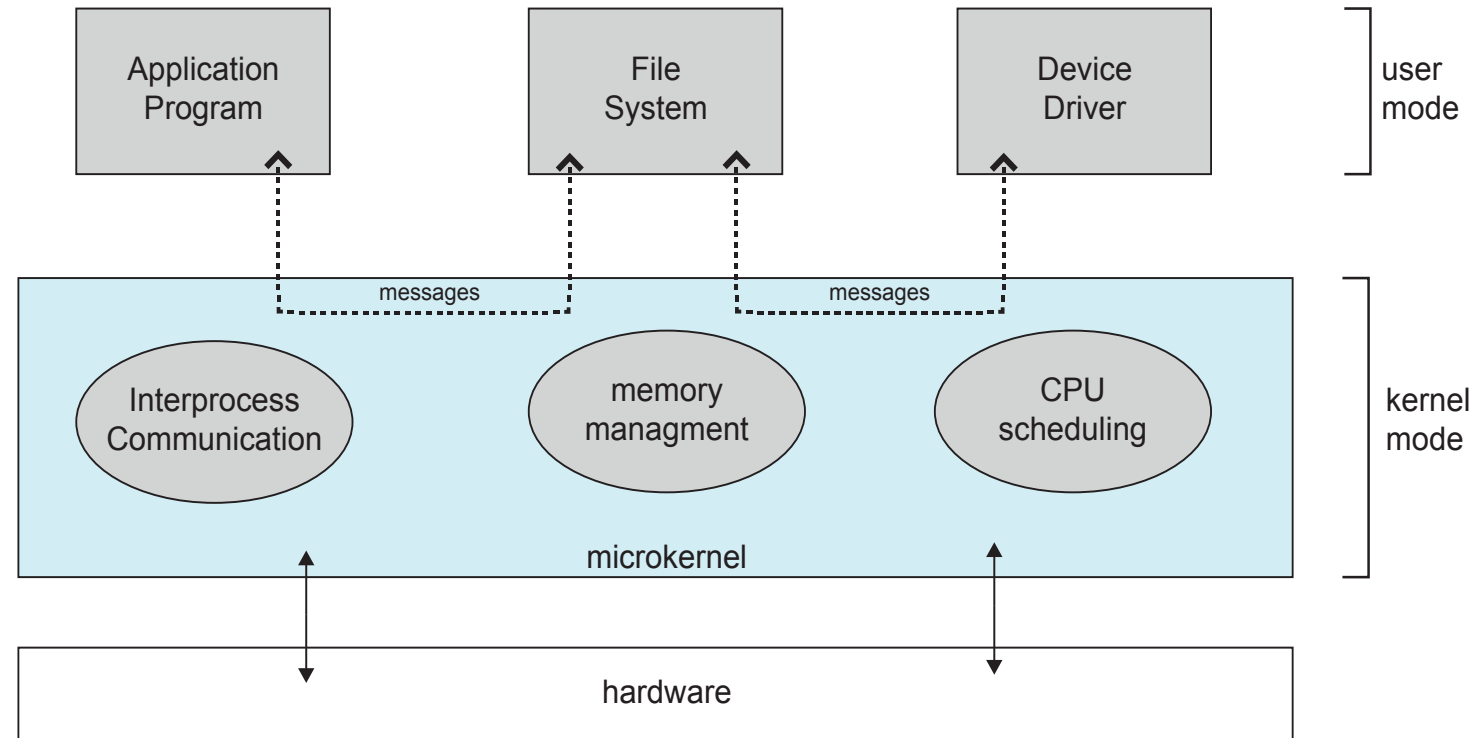
- The operating system is divided into a number of layers (levels), each built on top of lower layers
- The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers



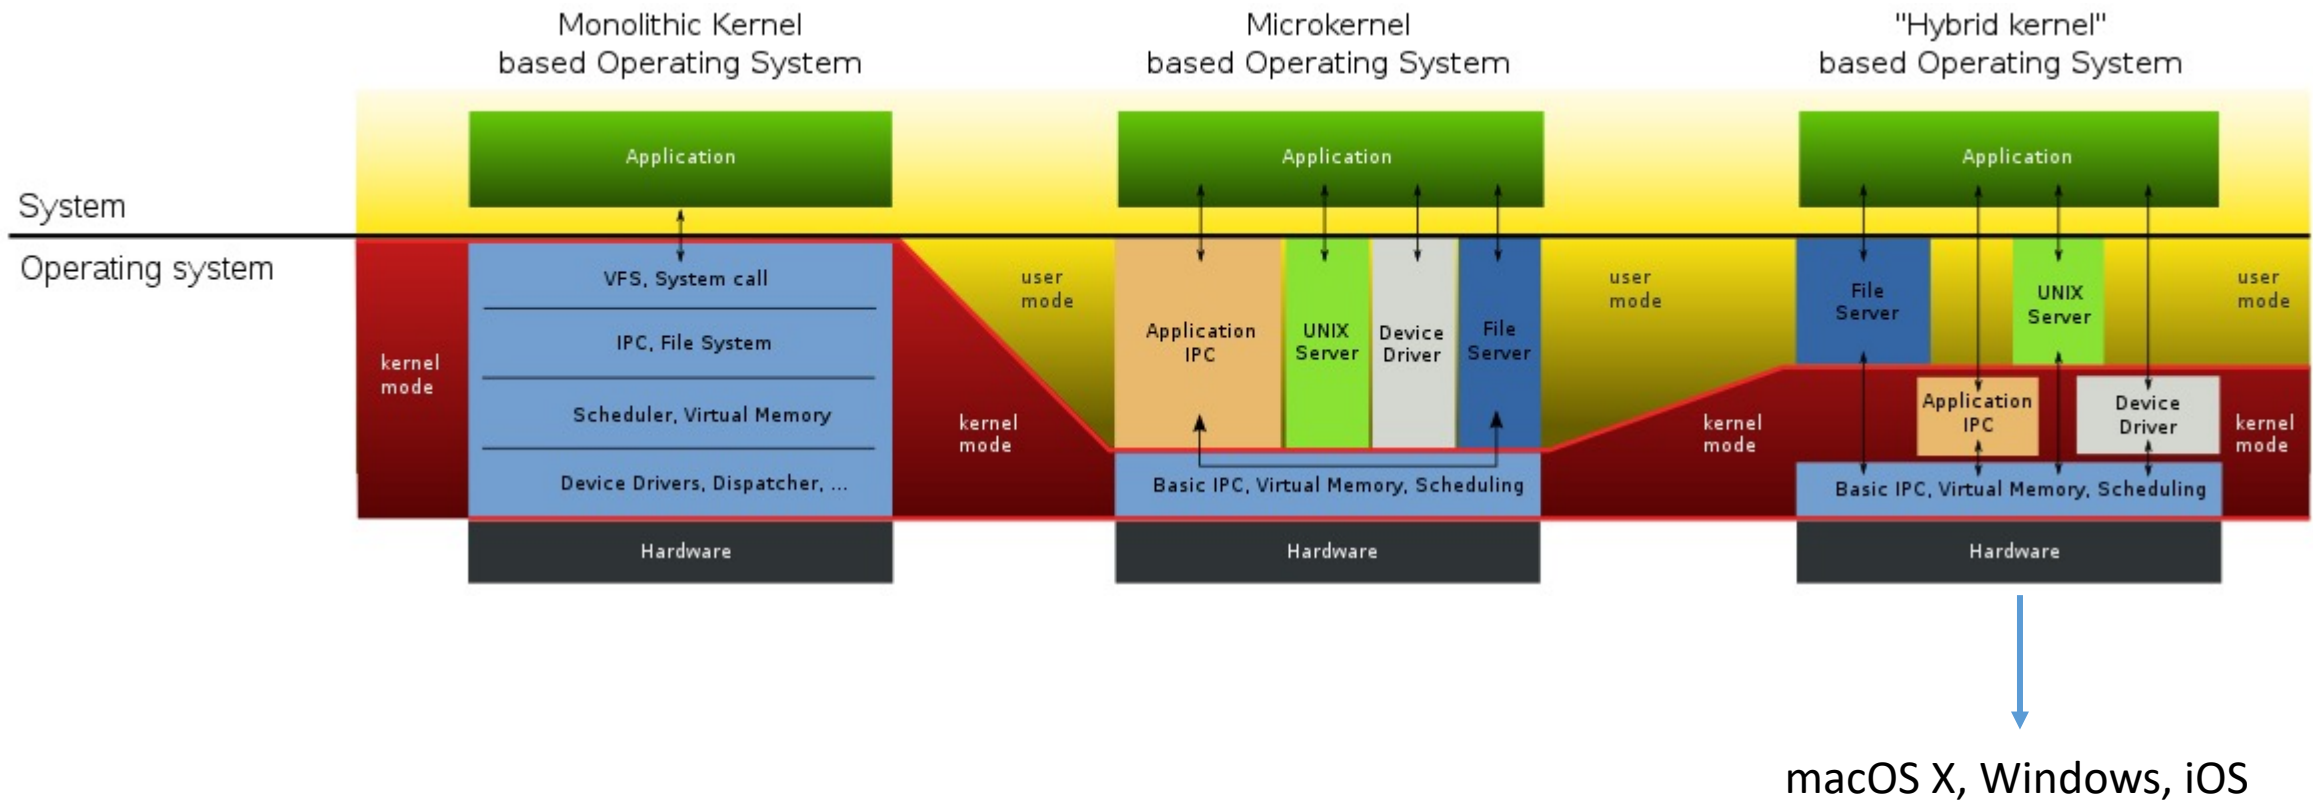
# Microkernel



- user services and kernel services are in separate address spaces
- smaller in size
- slower
- extendible, all new services are added to user space
- if a service crashes, working of microkernel is not affected
- more secure and reliable
- eg., Mach, QNX, Windows NT (initial release)
- Drawback ??? *Performance overhead of user space to kernel space communication*



# A Comparison



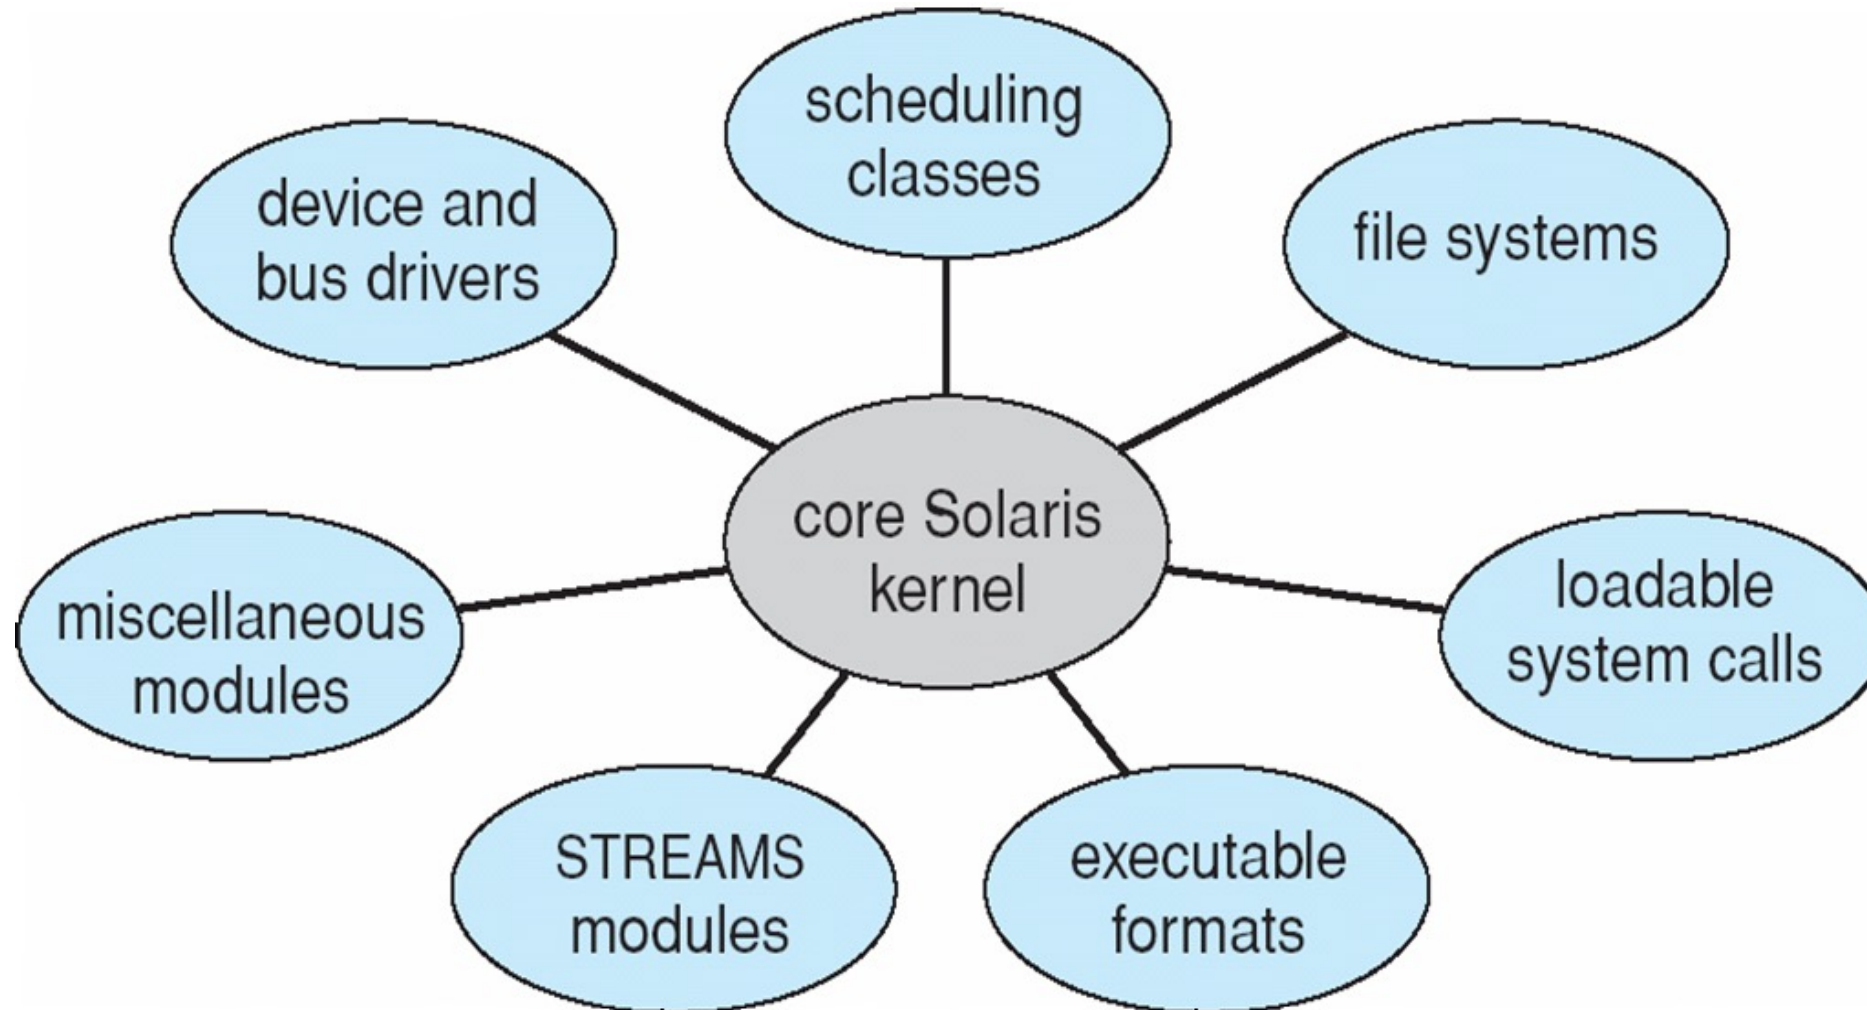


# Modules



- ❖ loadable kernel modules
- ❖ kernel has a core set of components
- ❖ links in additional services via modules, either at boot time or during run time
- ❖ each module has a well defined interface
- ❖ dynamically linking services is preferable to adding new features directly to the kernel → does not require recompiling the kernel for every change
- ❖ better than a layered approach → any module can call any module
- ❖ better than microkernel → no message passing required to invoke modules

# Modules



# Operating-System Debugging



Debugging is twice as hard as writing  
the code in the first place.  
Therefore, if you write the code as  
cleverly as possible, you are, by  
definition, not smart enough to  
debug it.

— *Brian Kernighan* —



# Performance Tuning


innovate

achieve

lead

```
top - 04:33:30 up 17:16, 1 user, load average: 0.05, 0.01, 0.03
Tasks: 75 total, 2 running, 73 sleeping, 0 stopped, 0 zombie
Cpu(s): 0.0%us, 0.3%sy, 0.0%ni, 88.6%id, 0.0%wa, 0.0%hi, 11.1%si, 0.0%st
Mem: 515348k total, 319956k used, 195392k free, 43432k buffers
Swap: 1048568k total, 0k used, 1048568k free, 201748k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
4	root	RT	-5	0	0	0	S	0.3	0.0	0:12.90	watchdog/0
32348	root	15	0	2200	996	796	R	0.3	0.2	0:00.04	top
1	root	15	0	2068	612	528	S	0.0	0.1	0:34.24	init
2	root	RT	-5	0	0	0	S	0.0	0.0	0:00.00	migration/0
3	root	34	19	0	0	0	S	0.0	0.0	0:00.04	ksoftirqd/0
5	root	10	-5	0	0	0	S	0.0	0.0	0:04.77	events/0
6	root	10	-5	0	0	0	S	0.0	0.0	0:01.84	khelper
7	root	11	-5	0	0	0	S	0.0	0.0	0:00.06	kthread
10	root	10	-5	0	0	0	S	0.0	0.0	0:00.72	kblockd/0
11	root	20	-5	0	0	0	S	0.0	0.0	0:00.00	kacpid
47	root	20	-5	0	0	0	S	0.0	0.0	0:00.00	cqueue/0
50	root	10	-5	0	0	0	S	0.0	0.0	0:00.00	khubd
52	root	10	-5	0	0	0	S	0.0	0.0	0:00.02	kseriod



# System Boot



- bootstrap program / bootstrap loader
- simple bootstrap loader fetches boot program from disk, which then starts the operating system
- program counter is loaded with the address of the first instruction in the boot program located in memory
- diagnostics to determine the status of the hardware
- *POST (Power-On Self-Test) is the sequence that a computer's basic system (or "starting program") checks the computer keyboard, random access memory, hard disk drives, and other hardware to ensure they are working correctly*

```
Diskette Drive B : None          Serial Port(s) : 3F0 2F0
Pri. Master Disk : LBA,ATA 100, 250GB Parallel Port(s) : 370
Pri. Slave Disk  : LBA,ATA 100, 250GB DDR at Bank(s) : 0 1 2
Sec. Master Disk : None
Sec. Slave Disk  : None

Pri. Master Disk HDD S.M.A.R.T. capability ... Disabled
Pri. Slave Disk  HDD S.M.A.R.T. capability ... Disabled

PCI Devices Listing ...
Bus  Dev  Fun  Vendor Device  SVID  SSID  Class  Device Class  IRQ
-----
0    27    0  8086  2668  1458  A005  0403  Multimedia Device  5
0    29    0  8086  2658  1458  2658  0C03  USB 1.1 Host Cntrlr  9
0    29    1  8086  2659  1458  2659  0C03  USB 1.1 Host Cntrlr 11
0    29    2  8086  265A  1458  265A  0C03  USB 1.1 Host Cntrlr 11
0    29    3  8086  265B  1458  265A  0C03  USB 1.1 Host Cntrlr  5
0    29    7  8086  265C  1458  5006  0C03  USB 1.1 Host Cntrlr  9
0    31    2  8086  2651  1458  2651  0101  IDE Cntrlr        14
0    31    3  8086  266A  1458  266A  0C05  SMBus Cntrlr      11
1     0    0  10DE  0421  10DE  0479  0300  Display Cntrlr    5
2     0    0  1283  8212  0000  0000  0180  Mass Storage Cntrlr 10
2     5    0  11AB  4320  1458  E000  0200  Network Cntrlr    12
                        ACPI Controller  9
```

# System Boot



```
Ubuntu 8.04, kernel 2.6.24-16-generic
Ubuntu 8.04, kernel 2.6.24-16-generic (recovery mode)
Ubuntu 8.04, memtest86+
```

Use the ↑ and ↓ keys to select which entry is highlighted.  
Press enter to boot the selected OS, 'e' to edit the  
commands before booting, or 'c' for a command-line.

---

# Thank You