



BITS Pilani
Hyderabad Campus

OPERATING SYSTEMS (CS F372)

Processes

Dr. Barsha Mitra
CSIS Dept., BITS Pilani, Hyderabad Campus

Process Concept


- ❖ An operating system executes a variety of programs
 - ❖ Batch system – **jobs**
 - ❖ Time-shared systems – **user programs** or **tasks**
- ❖ Terms ***job*** and ***process*** used interchangeably
- ❖ **Process** – a program in execution; process execution must progress in sequential fashion

Process Concept

- ❖ Multiple parts
 - ❖ The program code, also called **text section**
 - ❖ Current activity including **program counter**, processor registers
 - ❖ **Stack** containing temporary data
 - ❖ Function parameters, return addresses, local variables
 - ❖ **Data section** containing global variables
 - ❖ **Heap** containing memory dynamically allocated during run time

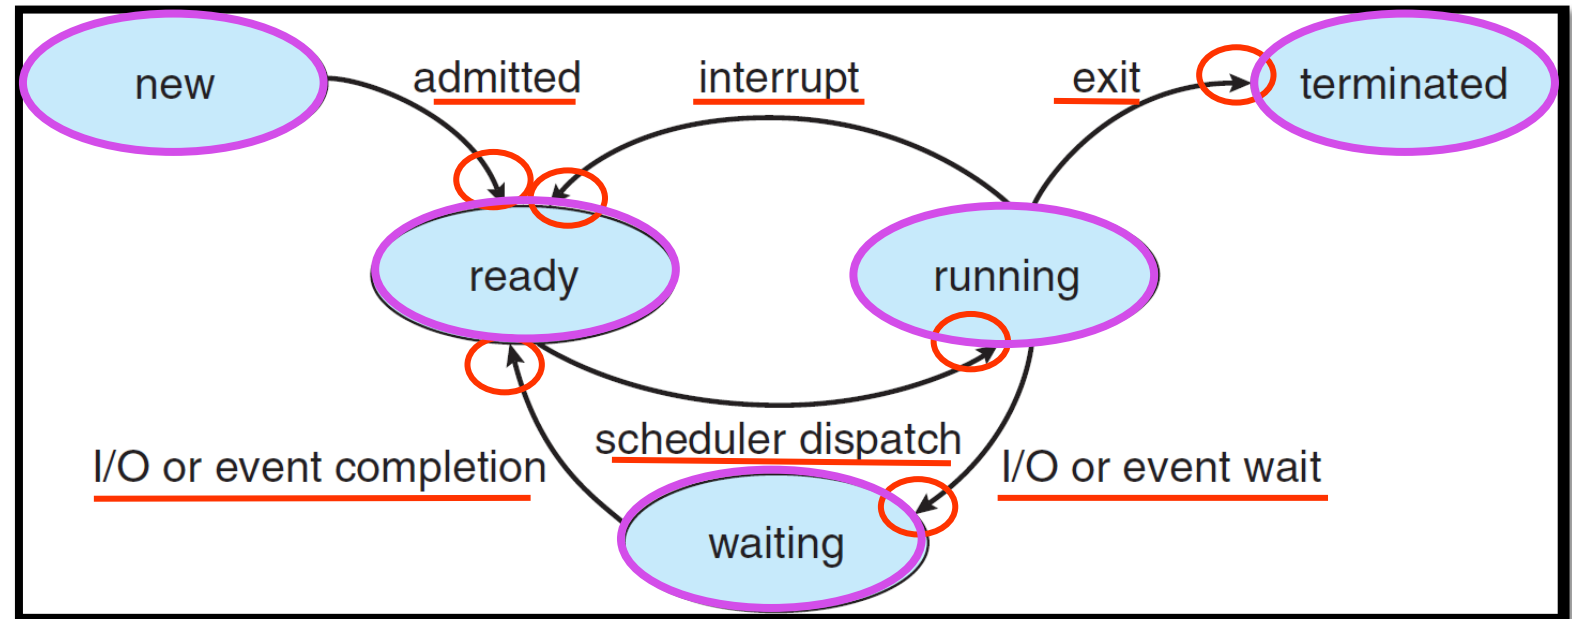
Process Concept

- ❖ Program is ***passive*** entity stored on disk (**executable file**), process is ***active***
 - ❖ Program becomes process when executable file loaded into memory
- ❖ Execution of program started via GUI mouse clicks, command line entry of its name, etc
- ❖ One program can be several processes
 - ❖ Consider multiple users executing the same program



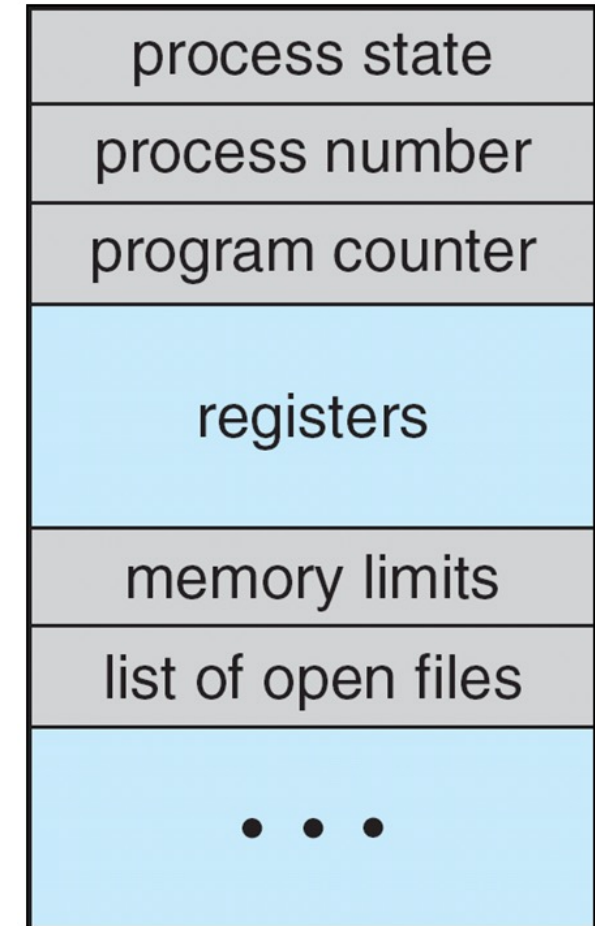
States of Process

- ❖ **new**: The process is being created
- ❖ **running**: Instructions are being executed
- ❖ **waiting**: The process is waiting for some event to occur
- ❖ **ready**: The process is waiting to be assigned to a processor
- ❖ **terminated**: The process has finished execution



Process Control Block

- ❖ **Process state** – new, ready, running, waiting, etc.
- ❖ **Program counter** – location of instruction to next execute
- ❖ **CPU registers** – contents of all process-centric registers
- ❖ **CPU scheduling information**- priorities, scheduling queue pointers, scheduling parameters
- ❖ **Memory-management information** – memory allocated to the process
- ❖ **Accounting information** – CPU used, clock time elapsed since start, time limits, account nos., process nos.
- ❖ **I/O status information** – I/O devices allocated to process, list of open files



Process Creation

- ❖ **Parent** process creates **children** processes, which, in turn create other processes, forming a **tree** of processes
- ❖ Generally, process identified and managed via a **process identifier (pid)**, integer number
- ❖ Resource sharing options (CPU time, memory, files, I/O devices)
 - ❖ Parent and children share all resources
 - ❖ Children share subset of parent's resources
 - ❖ Parent and child share no resources
- ❖ Execution options
 - ❖ Parent and children execute concurrently
 - ❖ Parent waits until children terminate

Process Creation

- ❖ Address space
 - ❖ Child duplicate of parent (same program and data)
 - ❖ Child has a program loaded into it
- ❖ UNIX examples
 - ❖ `fork()` system call creates new process
 - ❖ `exec()` system call used after a `fork()` to replace the process's memory space with a new program

Process Creation

❖ **fork()**

- ❖ address space of child process is a copy of parent process
- ❖ both child and parent continue execution at the instruction after fork()
- ❖ return code for fork() is 0 for child
- ❖ return code for fork() is non-zero (child *pid*) for parent

❖ **exec()**

- ❖ loads a binary file into memory and starts execution
- ❖ destroys previous memory image
- ❖ call to exec() does not return unless an error occurs

❖ **wait()**


- ❖ parent can issue wait() to move out of ready queue until the child is done

Process Creation

```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int main()
{
    pid_t pid;
    pid = fork(); /* fork a child process */
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        printf("Child Process\n");
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        wait(NULL); /* parent waits for child to complete */
        printf("Child Complete");
    }
    return 0;
}

```



OUTPUT:

Child Process

a.out

Documents examples.desktop MyPrograms Pictures Templates

Desktop Downloads Music parent.c Public Videos

Child Complete

Example 1



```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    fork();
    printf ("hello\n");
    return (0);
}
```

hello
hello

Example 2



```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int x;
    x = fork();
    if (x == 0 )
        printf ("Child Process :%d", x);
    else
        printf ("I am parent : %d", x);
    return (0);
}
```

Child Process : 0
I am Parent : 1234

Example 3



```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    fork();
    fork();
    fork();
    printf("Hello");
    return (0);
}
```

Process Creation

```
int main()
{
    pid_t pid;
    pid = fork();
    if (pid < 0) {
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) {
        printf("Child Process\n");
        printf("child pid = %d\n", getpid());
    }
    else {
        printf("parent pid = %d\n", getpid());
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

OUTPUT:

```
parent pid = 6597
Child Process
child pid = 6598
Child Complete
```



Process Creation

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int main()
{
    pid_t pid; int x = 10;
    pid = fork();
    if (pid < 0) {
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) {
        printf("Child Process: x = %d\n", x);
        execlp("/bin/ls", "ls", NULL);
    }
    else {
        printf("Parent Process: x = %d\n", x);
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

OUTPUT:

Parent Process: x = 10

Child Process: x = 10

a.out Documents examples.desktop MyPrograms Pictures Templates

Desktop Downloads Music parent.c Public Videos

Child Complete

Process Creation

```
int main()
{
    pid_t pid;
    int x = 10;
    pid = fork();
    if (pid < 0) {
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) {
        x = x + 10;
        printf("Child Process: x = %d\n", x);
    }
    else {
        wait(NULL);
        printf("Parent Process: x = %d\n", x);
        printf("Child Complete");
    }
    return 0;
}
```

OUTPUT:

Child Process: x = 20
Parent Process: x = 10
Child Complete

Process Creation

```
int main()
{
    pid_t pid;
    int x = 10;
    pid = fork();
    if (pid < 0) {
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) {
        for(long i = 0; i < 500000000000; i++);
        printf("Child Process: x = %d\n", x);
    }
    else {
        x = x + 10;
        wait(NULL);
        printf("Parent Process: x = %d\n", x);
        printf("Child Complete");
    }
    return 0;
}
```

OUTPUT:

Child Process: x = 10
Parent Process: x = 20
Child Complete

Process Termination

- ❖ Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
- ❖ May return status data from child to parent (via `wait()`)
- ❖ Process' resources are deallocated by operating system
- ❖ Parent may terminate the execution of children processes because:
 - ❖ Child has exceeded allocated resources limit
 - ❖ Task assigned to child is no longer required
 - ❖ Parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- ❖ Some operating systems do not allow child to exist if parent has terminated.
- ❖ If a process terminates, then all its children must also be terminated.
 - ❖ **cascading termination** - All children, grandchildren, etc. are terminated.
 - ❖ The termination is initiated by the operating system.
- ❖ The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process

```
pid_t pid;  
int status;  
pid = wait(&status); //parent can tell which child has terminated
```
- ❖ If no parent waiting (did not invoke **wait()** till then) process is a **zombie**
- ❖ If parent terminated, process is an **orphan**

Zombie Process

```
int main()
{
    pid_t child_pid = fork();
    if (child_pid > 0) {
        sleep(10); → zombie
        wait(NULL); → no longer a zombie
        sleep(200);
    }
    else{
        printf("\n%d",getpid());
        exit(0);
    }
    return 0;
}
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
1	Z	1001	17404	17403	0	80	0	-	0	-	pts/0	00:00:00	a.out <defunct>

Orphan Process

```
int main()
{
    pid_t child_pid = fork();
    if (child_pid > 0){
        printf("\nParent process: %d\n",getpid());
        sleep(6);
    }
    else{
        printf("\nParent PID: %d\n",getppid());
        sleep(20);
        printf("\nChild Process: %d",getpid());
        printf("\nParent PID: %d",getppid());
        exit(0);
    }

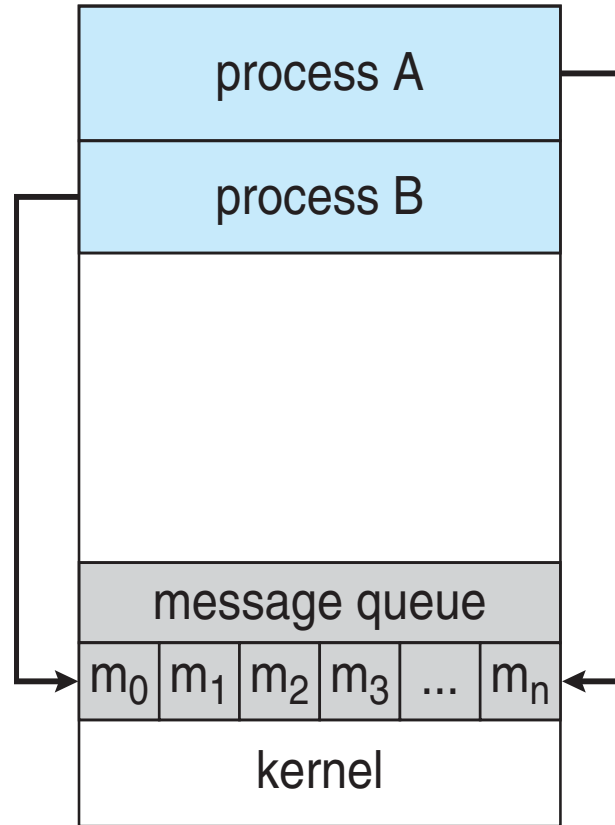
    return 0;
}
```

Interprocess Communication


- ❖ Processes within a system may be *independent* or *cooperating*
- ❖ Cooperating process can affect or be affected by other processes, including sharing data
- ❖ Reasons for cooperating processes:
 - ❖ Information sharing
 - ❖ Computation speedup
 - ❖ Modularity
 - ❖ Convenience
- ❖ Cooperating processes need **interprocess communication (IPC)**
- ❖ Two models of IPC
 - ❖ **Shared memory**
 - ❖ **Message passing**

Interprocess Communication

message passing



shared memory



Producer-Consumer Problem



IPC – Shared Memory

- ❖ An area of memory shared among the processes that wish to communicate
- ❖ The communication is under the control of the users processes not the operating system
- ❖ Provide mechanism that will allow the user processes to synchronize their actions when they access shared memory

Producer-Consumer Problem

- ❖ Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - ❖ **unbounded-buffer** places no practical limit on the size of the buffer
 - ❖ **bounded-buffer** assumes that there is a fixed buffer size

- ❖ Shared data, reside in a region of memory shared by producer & consumer


```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- ❖ Solution is correct, but can only use BUFFER_SIZE-1 elements

Bounded Buffer - Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out) //buffer is full
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```



```
item next_consumed;
while(true){
    while (in == out) //buffer is empty
        ; /*do nothing*/
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next consumed */
}
```

IPC – Message Passing

- ❖ Mechanism for processes to communicate and to synchronize their actions
- ❖ processes communicate with each other without resorting to shared variables, no sharing of address space
- ❖ IPC facility provides two operations:
 - ❖ **send(*message*)**
 - ❖ **receive(*message*)**
- ❖ The *message* size is either fixed or variable
- ❖ If processes *P* and *Q* wish to communicate, they need to:
 - ❖ Establish a **communication link** between them
 - ❖ Exchange messages via `send()` / `receive()`

Message Passing – Direct Communication



- ❖ Processes must name each other explicitly:
 - ❖ **send (*P*, *message*)** – send a message to process *P*
 - ❖ **receive(*Q*, *message*)** – receive a message from process *Q*
- ❖ Properties of communication link
 - ❖ Links are established automatically
 - ❖ Processes only need to know each other's identity
 - ❖ A link is associated with exactly one pair of communicating processes
 - ❖ Between each pair there exists exactly one link

Message Passing – Indirect Communication



- ❖ Messages are directed and received from mailboxes (also referred to as ports)
 - ❖ Each mailbox has a unique ID
 - ❖ Processes can communicate only if they share a mailbox
- ❖ Properties of communication link
 - ❖ Link is established only if processes share a common mailbox
 - ❖ A link may be associated with many processes
 - ❖ Each pair of processes may share several communication links, each link corresponds to one mailbox

Message Passing – Indirect Communication



- ❖ Operations
 - ❖ create a new mailbox (port)
 - ❖ send and receive messages through mailbox
 - ❖ destroy a mailbox
- ❖ Primitives are defined as:
 - ❖ **send(*A, message*)** – send a message to mailbox A
 - ❖ **receive(*A, message*)** – receive a message from mailbox A

Synchronization

- ❖ Message passing may be either blocking or non-blocking
- ❖ **Blocking** is considered **synchronous**
 - ❖ **Blocking send** -- the sender is blocked until the message is received by the receiving process or mailbox
 - ❖ **Blocking receive** -- the receiver is blocked until a message is available
- ❖ **Non-blocking** is considered **asynchronous**
 - ❖ **Non-blocking send** -- the sender sends the message and continues
 - ❖ **Non-blocking receive** -- the receiver receives:
 - ❖ A valid message, or
 - ❖ Null message

Buffering


- ❖ messages exchanged by communicating processes reside in a temporary queue
- ❖ implemented in one of three ways
 - ❖ **Zero capacity** – no messages are queued, link can't have any waiting messages, sender must block until receiver receives message
 - ❖ **Bounded capacity** – queue is of finite length of n messages, sender need not block if queue is not full, sender must wait if queue full
 - ❖ **Unbounded capacity** – infinite length queue, sender never blocks

Pipe

- ❖ Acts as a conduit allowing two processes to communicate
- ❖ Issues:
 - ❖ Is communication unidirectional or bidirectional?
 - ❖ In the case of two-way communication, is it half or full-duplex?
 - ❖ Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?
 - ❖ Can the pipes be used over a network?
- ❖ **Ordinary pipes** –
 - ❖ cannot be accessed from outside the process that created it
 - ❖ parent process creates a pipe and uses it to communicate with a child process that it created
- ❖ **Named pipes** – can be accessed without a parent-child relationship

Ordinary Pipe

- ❖ Ordinary Pipes allow communication in standard producer-consumer style
- ❖ Producer writes to one end (the **write-end** of the pipe)
- ❖ Consumer reads from the other end (the **read-end** of the pipe)
- ❖ Ordinary pipes are unidirectional
- ❖ Require parent-child relationship between communicating processes
- ❖ Windows calls these **anonymous pipes**



Ordinary Pipe

- ❖ ordinary pipe can't be accessed from outside the process that created it
- ❖ parent process creates a pipe and uses it to communicate with a child process that it creates via `fork()`
- ❖ child inherits the pipe from its parent process like any other file

Ordinary Pipe

```
#include <sys/types.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
}
```

```
    if (pid > 0) { /* parent process */
        /* close the unused end of the pipe */
        close(fd[READ_END]);

        /* write to the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1);

        /* close the write end of the pipe */
        close(fd[WRITE_END]);
    }

    else { /* child process */
        /* close the unused end of the pipe */
        close(fd[WRITE_END]);

        /* read from the pipe */
        read(fd[READ_END], read_msg, BUFFER_SIZE);
        printf("read %s", read_msg);

        /* close the write end of the pipe */
        close(fd[READ_END]);
    }


    return 0;
}
```

Named Pipe

- ❖ Named Pipes are more powerful than ordinary pipes
- ❖ Communication is bidirectional
- ❖ No parent-child relationship is necessary between the communicating processes
- ❖ Several processes can use the named pipe for communication
- ❖ Do not cease to exist if the communicating processes have terminated
- ❖ Provided on both UNIX and Windows systems
- ❖ Referred to as FIFOs in UNIX systems

Message Queue

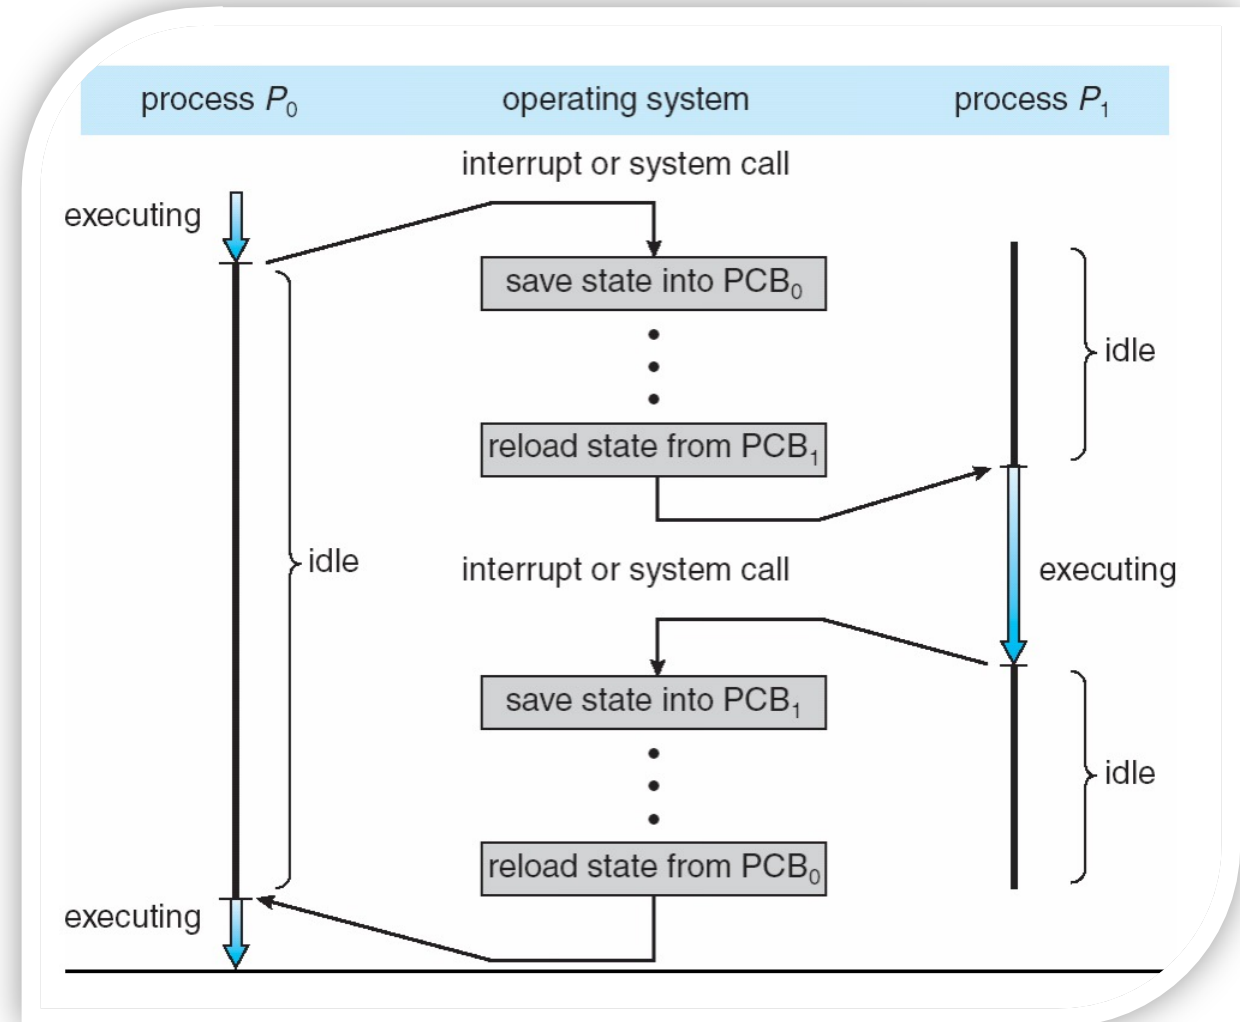
- asynchronous communication
- messages placed onto the queue are stored until the recipient retrieves them



- **Step 1** – Create a message queue or connect to an already existing message queue (`msgget()`)
- **Step 2** – Write into message queue (`msgsnd()`)
- **Step 3** – Read from the message queue (`msgrcv()`)
- **Step 4** – Perform control operations on the message queue (`msgctl()`)

Context Switch

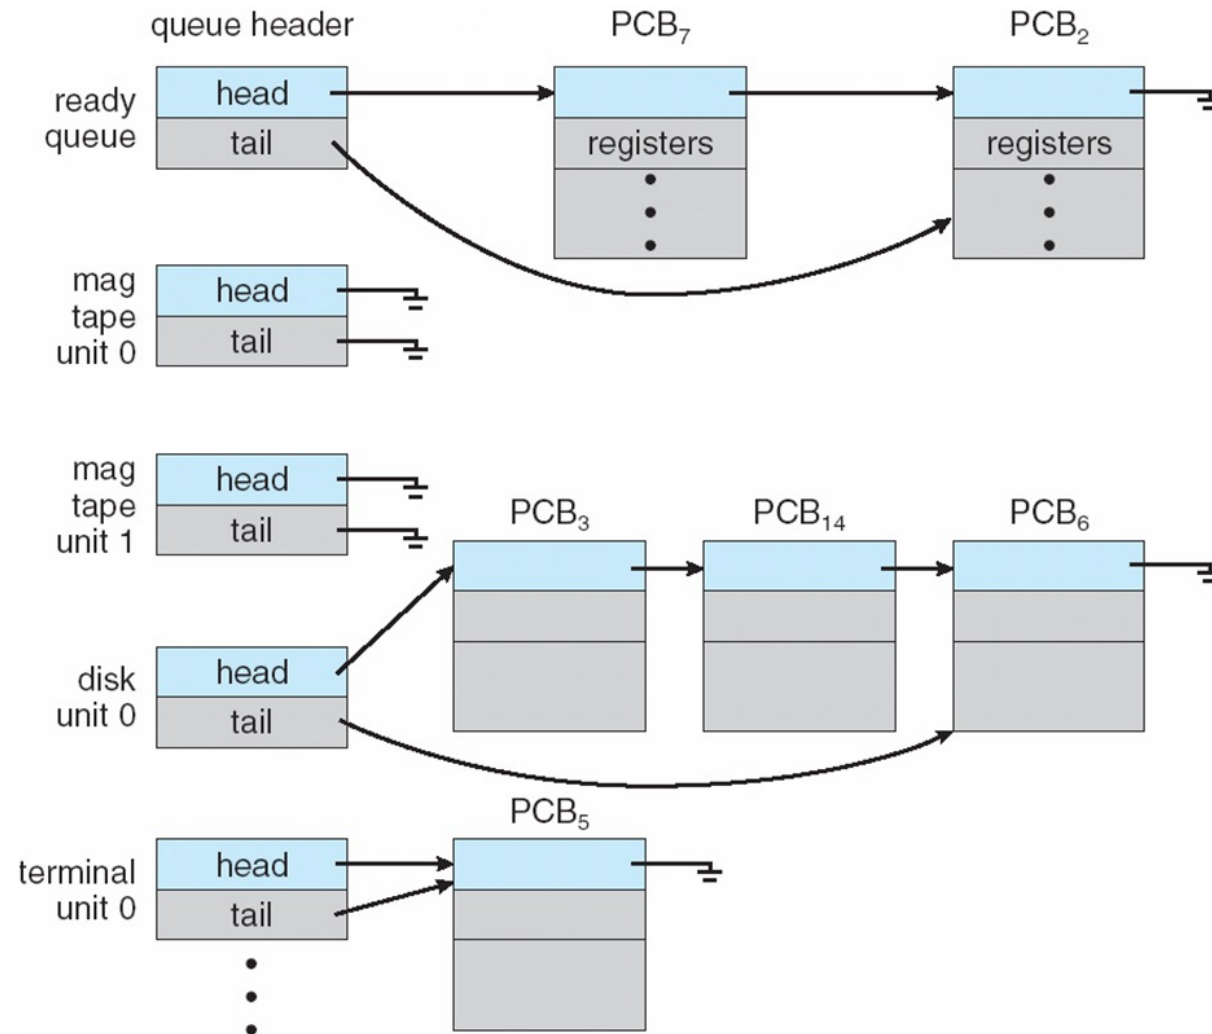
- ❖ When CPU switches to another process, system must **save state** of the old process and load the **state** for the new process via a **context switch**
- ❖ **Context** of a process represented in the PCB (CPU registers contents, process state, memory management info.)
- ❖ Context-switch time is overhead; system does no useful work while switching
- ❖ Time is dependent on hardware support



Process Scheduling

- ❖ Maximize CPU use, quickly switch processes onto CPU for time sharing
- ❖ **Process scheduler** selects among available processes for next execution on CPU
- ❖ Maintains **scheduling queues** of processes
 - ❖ **Job queue** – set of all processes in the system
 - ❖ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute, generally stored as a linked list
 - ❖ **Device queues** – set of processes waiting for an I/O device
- ❖ Processes migrate among the various queues

Various Queues

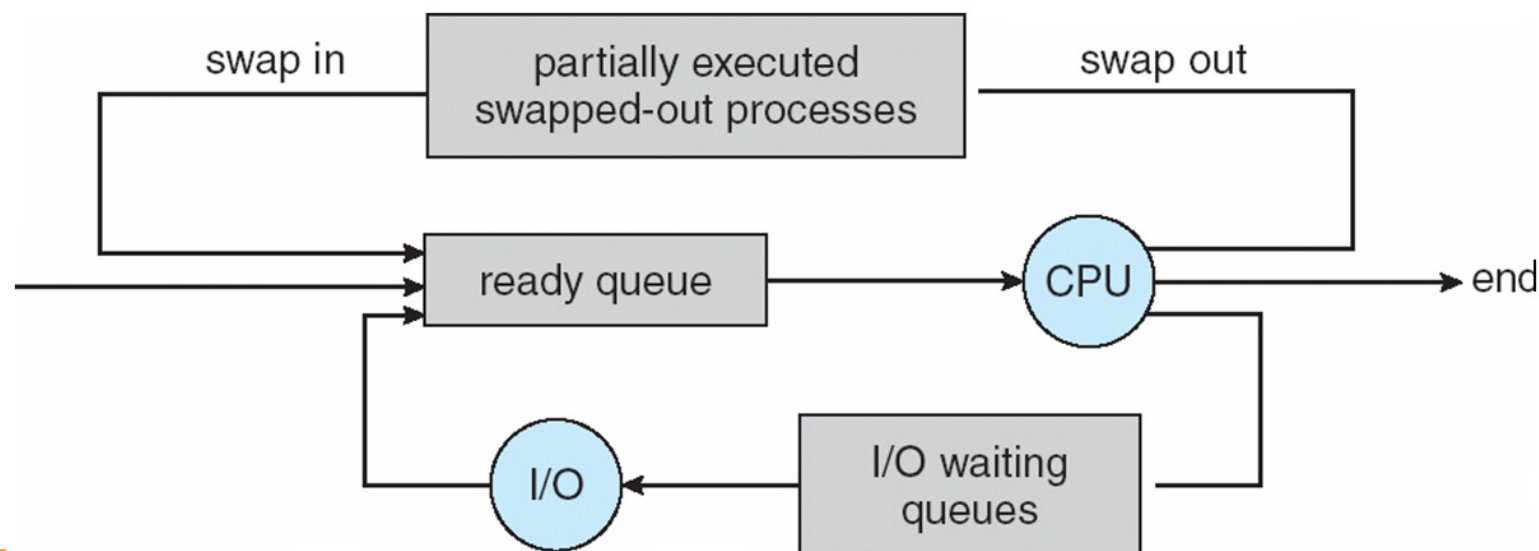


Schedulers

- ❖ **Short-term scheduler (or CPU scheduler)** – selects which process should be executed next and allocates CPU
 - ❖ Sometimes the only scheduler in a system
 - ❖ Short-term scheduler is invoked frequently (milliseconds) \Rightarrow (must be fast)
- ❖ **Long-term scheduler (or job scheduler)** – selects which processes from job queue should be brought into the ready queue
 - ❖ Long-term scheduler is invoked infrequently (seconds, minutes) \Rightarrow (may be slow)
 - ❖ The long-term scheduler controls the **degree of multiprogramming (number of processes in main memory)**
- ❖ Processes can be described as either:
 - ❖ **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - ❖ **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- ❖ Long-term scheduler strives for good ***process mix***

Schedulers

- ❖ **Medium-term scheduler** can be added in time sharing systems if degree of multiprogramming needs to decrease
 - ❖ Intermediate level of scheduling
 - ❖ Remove process from memory, store on disk, bring back in from disk to continue execution from where it left off: **swapping**
 - ❖ **Required for improving process mix or for freeing of memory**



Thank You